


Data Preprocessing: Cleaning and Organizing

Dr. Ratna Babu Chinnam
Industrial & Systems Engineering
Wayne State University

Motivation



What data scientists spend the most time doing

- Building training sets: 3%
- Cleaning and organizing data: 60%
- Collecting data sets: 19%
- Mining data for patterns: 9%
- Refining algorithms: 4%
- Other: 5%

According to a survey in Forbes (March 2016), data scientists spend **80%** of their time on **data preparation**. | [Link](#)

Handling Outliers*

CAUTION: Best way to detect outliers is to explore data visually. Numerical methods are prone to mistakes.

Outlier Detection with Standard Deviation (σ)

- If a value has a distance to the average higher than $k * \sigma$, it can be assumed as an outlier. What should k (factor) be?
 - No trivial solution for k , but usually, a value between (2, 4) seems practical.

- #Dropping outlier rows with σ

```
factor = 3
upper_lim = data['column'].mean() + data['column'].std() * factor
lower_lim = data['column'].mean() - data['column'].std() * factor
data = data[(data['column'] < upper_lim) & (data['column'] > lower_lim)]
```

Outlier Detection with Percentiles

- Another method to detect outliers is to use percentiles.
 - Assume a certain % of the value from the top or the bottom as an outlier.
 - Threshold depends on distribution of data.
 - If your data ranges from **0** to **100**, your top **5%** is not the values between **96** and **100**. Top **5%** means here the values that are out of the **95th** percentile of data.

- #Dropping the outlier rows with Percentiles

```
upper_lim = data['column'].quantile(.95)
lower_lim = data['column'].quantile(.05)
data = data[(data['column'] < upper_lim) & (data['column'] > lower_lim)]
```


*Python code assumes that you have imported **Pandas** and **Numpy** libraries.

An Outlier Dilemma: Drop or Cap

- If possible, detected outliers should be confirmed with domain experts
- One option for handling outliers is to **cap** them instead of dropping.
 - It might be better for final model performance.
 - Capping can affect the distribution of the data.
- #Capping the outlier rows with Percentiles
 - `upper_lim = data['column'].quantile(.95)`
 - `lower_lim = data['column'].quantile(.05)`
 - `data.loc[(df[column] > upper_lim), column] = upper_lim`
 - `data.loc[(df[column] < lower_lim), column] = lower_lim`

Variable Binning

- Process of converting a numerical variable into a categorical variable
 - Age: [0,5),[5,10),[10,20),[20,60),[60+)
 - Ensure that all bins have a decent frequency
- Binning can help deal with highly non-linear effects
 - May not be useful for ANNs; Might even degrade performance.




#Numerical Binning Example

- `data['bin'] = pd.cut(data['value'], bins=[0,30,70,100], labels=["Low", "Mid", "High"])`

	value	bin
0	2	Low
1	45	Mid
2	7	Low
3	85	High
4	28	Low

Logarithm Transformation

- Log transform is commonly used in feature engineering to deal with skewed data.
 - May not be useful for ANNs; Might even degrade performance.
 - It also decreases the effect of the outliers, due to the normalization of magnitude differences and the model become more robust.
- Data must have only + values, otherwise you receive an error.
 - Add a constant to your data before transformation to avoid issues.



Source: McDonald | [Link](#)

#Log Transform Example

- `data = pd.DataFrame({'value': [2, 45, -23, 85, 28, 2, 35, -12]})`
 - `data['log+1'] = (data['value']+1).transform(np.log)`
 - #Negative Values Handling
 - #Note that the values are different
 - `data['log'] = (data['value'] - data['value'].min()+1).transform(np.log)`
- | | value | $\log(x+1)$ | $\log(x - \min(x) + 1)$ |
|---|-------|-------------|-------------------------|
| 0 | 2 | 1.09861 | 3.25810 |
| 1 | 45 | 3.82864 | 4.23411 |
| 2 | -23 | nan | 0.00000 |
| 3 | 85 | 4.45435 | 4.69135 |
| 4 | 28 | 3.36730 | 3.95124 |
| 5 | 2 | 1.09861 | 3.25810 |
| 6 | 35 | 3.58352 | 4.07754 |
| 7 | -12 | nan | 2.48491 |

One-Hot Encoding

- Spreads the values in a column to multiple flag columns and assigns 0 or 1 to them.
 - These binary values express the relationship between grouped and encoded column.
- If you have N distinct values in the column, it is enough to map them to N-1 binary columns, because the missing value can be deducted from other columns.
- `encoded_columns =
pd.get_dummies(data['column'])`
- `data =
data.join(encoded_columns).drop('column',
axis=1)`

Scaling

- In many cases, numerical features of dataset do not have a certain range and they differ from each other.
- Scaling data to a fixed range (e.g., [0,1]) or standardizing it by making the $\mu = 0$ and $\sigma = 1$ can help accelerate learning.

- **Fixed Range:**

- `data = pd.DataFrame({'value': [2, 45, -23, 85, 28, 2]})`

- `data['normalized'] = (data['value'] - data['value'].min()) / (data['value'].max() - data['value'].min())`

	value	normalized
0	2	0.23
1	45	0.63
2	-23	0.00
3	85	1.00

4	28	0.47
5	2	0.23

- **Standardizing:**

- `data = pd.DataFrame({'value': [2, 45, -23, 85, 28, 2]})`

- `data['standardized'] = (data['value'] - data['value'].mean()) / data['value'].std()`

	value	standardized
0	2	-0.52
1	45	0.70
2	-23	-1.23
3	85	1.84
4	28	0.22
5	2	-0.52

Extracting Date

- Date/time columns might provide valuable information about the model target (e.g., sales on weekdays vs weekends; ER admissions by hour of day)
- Three types of preprocessing for dates:
 - Extracting the parts of the date into different columns: Year, month, day, etc.
 - Extracting the time period between the current date and columns in terms of years, months, days, etc.
 - Extracting some specific features from the date: Name of the weekday, Weekend or not, holiday or not, etc.
- `from datetime import date`
- `data = pd.DataFrame({'date': ['01-01-2017', '04-12-2008', '23-06-1988', '25-08-1999', '20-02-1993',]})`
- **#Transform string to date**

```
data['date'] = pd.to_datetime(data.date,
                              format="%d-%m-%Y")
```
- **#Extracting Year**

```
data['year'] = data['date'].dt.year
```
- **#Extracting Month**

```
data['month'] = data['date'].dt.month
```
- **#Extracting passed years since the date**

```
data['passed_years'] = date.today().year -
data['date'].dt.year
```

- **#Extracting passed months since the date**

```
data['passed_months'] = (date.today().year -
data['date'].dt.year) * 12 +
date.today().month - data['date'].dt.month
```

- **#Extracting the weekday name of the date**

```
data['day_name'] = data['date'].dt.day_name()
```

	date	year	month	passed_years	passed_months	day_name
0	2017-01-01	2017	1	2	26	Sunday
1	2008-12-04	2008	12	11	123	Thursday
2	1988-06-23	1988	6	31	369	Thursday
3	1999-08-25	1999	8	20	235	Wednesday
4	1993-02-20	1993	2	26	313	Saturday

References

- Jacques Peeters (April 2020) A framework for feature engineering and machine learning pipelines | [Link](#)
 - Very good general outline on how to write effective code
- Emre Rençberoğlu (April 2019) Fundamental Techniques of Feature Engineering for Machine Learning | [Link](#)
- Ways to Detect and Remove the Outliers | [Link](#)
- Understanding Feature Engineering:
 - Continuous Numeric Data | [Link](#)
 - Categorical Data | [Link](#)
- Log Transformations for Skewed and Wide Distributions | [Link](#)
- Tidy data | [Link](#)
- About Feature Scaling and Normalization | [Link](#)